



RISC-V Boot and Runtime Services Specification (BRS)

OS-A SEE Task Group

Version v0.0.1, 2024-04-15: This document is in development. Assume everything can change. See
<http://riscv.org/spec-state> for details.

Table of Contents

Preamble	1
Copyright and license information	2
Contributors	3
1. Introduction	4
1.1. Releases	4
1.2. Approach to Solutions	4
1.3. Testing and Conformance	4
1.4. Glossary	5
2. Recipes	7
2.1. BRS-I Recipe	7
2.2. BRS-B Recipe	7
3. Hart Requirements	9
4. SBI Requirements	10
5. BRS-I UEFI Requirements	11
5.1. BRS-I Security Requirements	11
5.2. BRS-I I/O-specific Requirements	12
5.3. BRS-I UEFI Runtime Services	12
5.4. BRS-I Firmware Update	13
6. BRS-I ACPI Requirements	14
6.1. BRS-I ACPI Methods and Objects	15
6.2. RVI-specific ACPI IDs	16
6.3. RVI-specific ACPI Device Properties	16
6.3.1. Properties for UART Devices	17
7. BRS-I SMBIOS Requirements	18
7.1. Type 04 Processor Information	19
7.2. Type 44 Processor-Specific Data	19
7.3. Processor-Specific Data Structure Versioning	20
8. Firmware Implementation Guidance	21
8.1. Recipes Guidance	21
8.1.1. BRS-I Recipe Guidance	21
8.2. UEFI Implementation Guidance	21
8.2.1. Privilege Levels	21
8.2.2. Firmware Update	21
8.2.3. PCIe	22
8.3. ACPI Implementation Guidance	22
8.3.1. 64-bits Clean	22
8.3.2. Hardware-Reduced ACPI	22
8.3.3. Table Guidance	22

8.3.4. DSDT and SSDTs	23
8.3.5. FADT	24
8.3.6. MADT	24
8.3.7. PLIC/APLIC Namespace Devices	24
8.3.8. PCIe	25
8.3.9. SPCR	25
8.4. SMBIOS Implementation Guidance	25
8.4.1. Type 44 Processor-Specific Data	25
Bibliography	26

Preamble



This document is in the [Development state](#)

Assume everything can change. This draft specification will change before being accepted as standard, so systems made to this draft specification will likely not conform to the future standard.

Copyright and license information

This specification is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at creativecommons.org/licenses/by/4.0/.

Copyright 2024 by RISC-V International.

Contributors

This RISC-V specification has been contributed to directly or indirectly by (in alphabetical order):

Aaron Durbin, Andrei Warkentin, Andrew Jones, Atish Patra, Darius Rad, Heinrich Schuchardt,
Jared McNeill, Paul Walmsley, Sunil V L, Vedvyas Shanbhogue

Chapter 1. Introduction

The *RISC-V Boot and Runtime Services Specification* (BRS) defines a standardized set of software capabilities, that portable system software, such as operating systems and hypervisors, can rely on being present in an implementation to utilize in acts of device discovery, OS boot and hand-off, system management, and other operations.

The BRS specification is targeting systems that implement S/U privilege modes, and optionally the HS privilege mode. This is the expected deployment for OSVs and system vendors in a typical ecosystem covering client systems up through server systems where software is provided by different vendors than the system vendor.

This specification standardizes the requires for software interfaces and capabilities by building on top of relevant industry and ratified RISC-V standards.

1.1. Releases

It is expected that the BRS will periodically release a new specification. The determination of a new release will be based on the evaluation of significant changes to its underlying dependencies.

1.2. Approach to Solutions

The BRS focuses on two solutions in the form of what is deemed a recipe. Each recipe contains the requirements needed to fulfill each solution. The requirements of each recipe will be marked accordingly with a unique identifier. The recipes are BRS-I (Interoperable) and BRS-B (Bespoke).

1.3. Testing and Conformance

To be compliant with this specification, an implementation **MUST** support all mandatory requirements and **MUST** support the listed versions of the specifications. This standard set of capabilities **MAY** be extended by a specific implementation with additional standard on custom capabilities, including compatible later versions of listed standard specifications. Portable system software **MUST** support the specified mandatory capabilities to be compliant with this specification.

The requirements in this specification use the following format:

ID#	Requirement
CAT_NNN	<p>The CAT is a category prefix that logically groups the requirements and is followed by 3 digits - NNN - assigning a numeric ID to the requirement.</p> <p>The requirements use the key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "NOT RECOMMENDED", "MAY", and "OPTIONAL" that are to be interpreted as described in RFC 2119 [1] when, and only when, they appear in all capitals, as shown here. When these words are not capitalized, they have their normal English meanings.</p>
<p><i>A requirement or a group of requirements may be followed by non-normative text providing context or justification for the requirement. The non-normative text may also be used to reference sources that are the origin of the requirement.</i></p>	

1.4. Glossary

Most terminology has the standard RISC-V meaning. This table captures other terms used in the document. Terms in the document prefixed by **PCIe** have the meaning defined in the *PCI Express Base Specification* [2] (even if they are not in this table).

Table 1. Terms and definitions

Term	Definition
ACPI	<i>Advanced Configuration and Power Interface Specification</i> [3].
BRS	<i>RISC-V Boot and Runtime Services Specification</i> . This document.
BRS-I	Boot and Runtime Services recipe targeting interoperation across different software suppliers.
BRS-B	Boot and Runtime Services recipe using a bespoke solution.
DT	DeviceTree [4].
EBBR	<i>Embedded Base Boot Requirements Specification</i> [5].
OSV	Operating System Vendor.
OS	Operating System or Hypervisor.
Profile	RISC-V Profile [6].
SBI	<i>RISC-V Supervisor Binary Interface Specification</i> [7].
SMBIOS	<i>System Management BIOS (SMBIOS) Reference Specification</i> [8].
SoC	System on a chip, a combination of processor and supporting chipset logic in single package.
System	A system is the entirety of a computing entity, including all hardware, firmware and software (hypervisor, operating system, user applications, user data). A system can be thought of both as a logical construct (e.g. a software stack) or physical construct (e.g. a notebook, a desktop, a server, a network switch, etc).

Term	Definition
UEFI	<i>Unified Extensible Firmware Interface Specification</i> [9] .

Chapter 2. Recipes

In this context, a recipe is a collection of firmware specification requirements that hardware, firmware, and software providers can implement to increase the likelihood that software written to the recipe will run predictably on all conforming devices.

The BRS specification defines two recipes: BRS-I (for "Interoperable") and BRS-B (for "Bespoke").

2.1. BRS-I Recipe

The BRS-I recipe aims to simplify end-user experiences, software compatibility and OS distribution, by defining a common specification for boot and runtime interfaces. BRS-I is expected to be used by general-purpose compute devices such as servers, desktops, laptops and other devices with industry expectations on silicon vendor, OS and software ecosystem interoperability. BRS-I enables operating system providers to build a single **generic** operating system image that can be **successfully booted** on compliant systems. **Generic** means not requiring system-specific customizations - only an implementation of BRS-I requirements. **Successfully boot** means basic system configuration, sufficient for detecting the need for system-specific drivers and loading such drivers.

It is understood that systems will deliver features beyond those covered by BRS-I. However, software written against a specific version of BRS-I must run, unaltered, without **anomalous and unexpected behavior** on systems that include such features and that are compliant to that specific version of BRS-I. Such behavior, caused by factors entirely unknown to a generic OS, is hard to diagnose and always results in a terrible user experience that negatively affects the value of the whole RISC-V standards-based ecosystem. **Anomalous and unexpected behavior** is taken to mean system instability and worst-case behavior for non-specialized workloads, but does not include suboptimal/unoptimized behavior or missing I/O or accelerator drivers.

A key tenet of BRS-I is constraining behavior outside the scope of BRS-I. Features violating the principle of least surprise and causing anomalous and unexpected behavior in a generic OS must be configured by firmware as opt-in. [See additional guidance](#).

Table 2. BRS-I Recipe Overview

Profile	UEFI	ACPI	DT	SBI	SMBIOS
>= RVA20	>= 2.10	>= 6.6	N/A	>= 2.0	>= 3.7.0

2.2. BRS-B Recipe

BRS-B is intended for cases where only a minimal level of firmware interaction is mandated, focusing primarily on the boot process. The BRS-B recipe is the simpler of the two BRS recipes. It is expected to be used by software that is tailored to specific devices. Examples include many types of mobile devices, devices with real time response requirements, or embedded devices running rich operating systems with custom distributions.

Table 3. BRS-B Recipe Overview

Profile	UEFI	ACPI	DT	SBI	SMBIOS
>= RVA20	EBBR, >= 2.1.0 [5]	optional, >= 6.6	optional, >= v0.3	>= 2.0	optional, >= 3.7.0

Either of ACPI table or Device Tree must be supplied, but never both at the same time.

Chapter 3. Hart Requirements

A compliant system includes a RISC-V application processor and the requirements in this section apply solely to harts in the application processors of a system.

The BRS specification is minimally prescriptive on the RISC-V hart requirements. It is anticipated that detailed requirements will be driven by target market segment and product/solution requirements.

ID#	Requirement
HR_010	The RISC-V application processor harts MUST be compliant to RVA20S64 profile [6].
<i>The BRS governs the interactions between 64-bit OS supervisor-mode software and 64-bit firmware. These are minimum requirements allowing for the wide variety of existing and future hart implementations to be supported. It is expected that operating systems and hypervisors may impose additional profile/ISA requirements, depending on the use-case and application.</i>	

Chapter 4. SBI Requirements

The *RISC-V Supervisor Binary Interface Specification* (SBI) [7] defines the interface between the supervisor mode and the next higher privilege mode. This section defines the mandatory SBI version and extensions implemented by the higher privilege mode in order to be compatible with this specification.

ID#	Requirement
SBI_010	The SBI implementation MUST conform to SBI v2.0 or later.
SBI_020	The SBI implementation MUST implement the Hart State Management (HSM) extension.
<i>HSM is used by an OS for starting up, stopping, suspending and querying the status of secondary harts.</i>	

Certain requirements are conditional on the presence of RISC-V ISA extensions or system features.

ID#	Requirement
SBI_030	The Timer Extension (TIME) MUST be implemented, if the RISC-V "stimecmp / vstimecmp" Extension (Sstc) [10] is not available.
SBI_040	The S-Mode IPI Extension MUST be implemented, if the Incoming MSI Controller (IMSIC) [11] is not available.
SBI_050	The RFENCE Extension (RFNC, [7]) extension MUST be implemented, if the Incoming MSI Controller (IMSIC) [11] is not available.
SBI_060	The Performance Monitoring Extension (PMU) MUST be implemented, if the counter delegation-related ISA extensions (S*csrind [12], Smcdeleg [13], Sccfg [13]) are not present.
<i>NOTE: The PMU extension is currently being developed by the performance analysis TG [14].</i>	

Chapter 5. BRS-I UEFI Requirements

The *Unified Extensible Firmware Interface Specification* (UEFI) describes the interface between the OS and the supervisor-mode firmware.

This section defines the BRS-I mandatory and optional UEFI requirements on top of existing [9] specification requirements. Additional non-normative guidance may be found in the [firmware implementation guidance](#) section.



All content in this section is optional and recommended for BRS-B.

ID#	Requirement
UEFI_010	Implement a 64-bit UEFI firmware.
UEFI_020	Meet the 3rd Party UEFI Certificate Authority (CA) requirements on UEFI memory mitigations [15].
UEFI_030	Meet BRS-I specific memory map requirements: <ul style="list-style-type: none">• The default memory space attribute MUST be <code>EFI_MEMORY_WB</code>.• Enable address translation.• Only use <code>EfiRuntimeServicesData</code> memory type for describing any SMBIOS data structures.
UEFI_040	An implementation MAY comply with the <i>UEFI Platform Initialization Specification</i> [16].
UEFI_050	All hart manipulation internal to a firmware implementation SHOULD be done before completion of the <code>EFI_EVENT_GROUP_READY_TO_BOOT</code> event, with all secondary harts remaining offline from that point on.
<i>This ensures an OS loader is entered with an OS-compatible state for all harts.</i>	
UEFI_060	Declare the <code>EFI_CONFORMANCE_PROFILES_UEFI_SPEC_GUID</code> conformance profile.
<i>The <code>EFI_CONFORMANCE_PROFILES_UEFI_SPEC_GUID</code> conformance profile MUST be declared, as the BRS requirements are a superset of UEFI [9] (Section 2.6).</i>	
UEFI_070	Declare the <code>EFI_CONFORMANCE_PROFILE_BRS_1_0_SPEC_GUID</code> conformance profile (<code>{ 0x05453310, 0x0545, 0x0545, { 0x05, 0x45, 0x33, 0x05, 0x45, 0x33, 0x05, 0x45 }}</code>).
<i>Only a system fully compliant to the requirements in this section MAY declare the <code>EFI_CONFORMANCE_PROFILE_BRS_1_0_SPEC_GUID</code> conformance profile.</i>	

5.1. BRS-I Security Requirements

ID#	Requirement
USEC_010	Systems implementing UEFI secure boot are RECOMMENDED to implement the <code>EFI_SECURITY_ARCH_PROTOCOL</code> and <code>EFI_SECURITY2_ARCH_PROTOCOL</code> protocols [16].

ID#	Requirement
<i>The Security and Security2 architectural protocols are overridden by some bootloaders (e.g. systemd-boot) to validate EFI binaries that cannot be validated against the UEFI security database.</i>	
USEC_020	Systems implementing a TPM MUST implement the <i>TCG EFI Protocol Specification</i> [17].

See additional [requirements for UEFI runtime services](#).

5.2. BRS-I I/O-specific Requirements

ID#	Requirement
UIO_010	Systems implementing PCIe MUST always initialize all root complex hardware and perform resource assignment for all endpoints and usable hotplug-capable switches in the system, even in a boot scenario from a non-PCIe boot device.
<i>This is a stronger requirement than the PCI Firmware Specification firmware/OS device hand-off state ([18] Section 3.5). See additional guidance.</i>	
UIO_020	Systems implementing <code>EFI_GRAPHICS_OUTPUT_PROTOCOL</code> SHOULD configure the frame buffer to be directly accessible.
<i>That is, <code>EFI_GRAPHICS_PIXEL_FORMAT</code> is not <code>PixelBltOnly</code> and <code>FrameBufferSize</code> is reported as a valid hart memory-mapped I/O address.</i>	

5.3. BRS-I UEFI Runtime Services

ID#	Requirement
URT_010	Systems without a Real-Time Clock (RTC) MUST meet the following requirements: <ul style="list-style-type: none"> <code>GetTime()</code> MUST be implemented (e.g. in terms of CPU cycle counter). <code>SetTime()</code> MUST return <code>EFI_UNSUPPORTED</code>, and be appropriately described in the <code>EFI_RT_PROPERTIES_TABLE</code>.
URT_020	Systems with a Real-Time Clock on an OS-managed bus (e.g. I2C, subject to arbitration issues due to access to the bus by the OS) MUST meet the following requirements: <ul style="list-style-type: none"> <code>GetTime()</code> and <code>SetTime()</code> MUST return <code>EFI_UNSUPPORTED</code>, when called after the UEFI boot services have been exited. <code>GetTime()</code> and <code>SetTime()</code> MUST be appropriately described in the <code>EFI_RT_PROPERTIES_TABLE</code>.
<i>Also see AML_070.</i>	
URT_030	The UEFI <code>ResetSystem()</code> runtime service MUST be implemented.

ID#	Requirement
	<i>The OS MUST call the <code>ResetSystem()</code> runtime service call to reset or shutdown the system, preferring this to SBI, ACPI or other system-specific mechanisms. This allows for systems to perform any required system tasks on the way out (e.g. servicing <code>UpdateCapsule()</code> or persisting non-volatile variables in some systems).</i>
URT_040	Non-volatile UEFI variables MUST persist across calls to the <code>ResetSystem()</code> runtime service call.
URT_050	UEFI runtime services MUST be able to update the variables directly without the aid of an OS.
URT_060	<p>The following requirements MUST be met for systems with UEFI secure boot:</p> <ul style="list-style-type: none"> • MUST support a minimum of 128 KiB of non-volatile storage for UEFI variables. • The maximum supported variable size MUST be at least 64 KiB. • The <code>db</code> signature database variable (<code>EFI_IMAGE_SECURITY_DATABASE</code>) MUST be created with <code>EFI_VARIABLE_TIME_BASED_AUTHENTICATED_WRITE_ACCESS</code>, to prevent rollback attacks. • The <code>dbx</code> signature database variable (<code>EFI_IMAGE_SECURITY_DATABASE1</code>) MUST be created with <code>EFI_VARIABLE_TIME_BASED_AUTHENTICATED_WRITE_ACCESS</code>, to prevent rollback attacks.

5.4. BRS-I Firmware Update

ID#	Requirement
UFU_010	Systems with in-band firmware updates MUST do so either via <code>UpdateCapsule()</code> UEFI runtime service ([9] Section 8.5.3) or via <i>Delivery of Capsules via file on Mass Storage Device</i> ([9] Section 8.5.5).
	<i>In-band means the firmware running on a hart updates itself.</i>
UFU_020	Systems implementing in-band firmware updates via <code>UpdateCapsule()</code> MUST accept updates in the <i>Firmware Management Protocol Data Capsule Structure</i> format as described in <i>Delivering Capsules Containing Updates to Firmware Management Protocol</i> [9] (Section 23.3).
UFU_030	Systems implementing in-band firmware updates via <code>UpdateCapsule()</code> MUST provide an ESRT [9] (Section 23.4) describing every firmware image that is updated in-band.
UFU_040	Systems implementing in-band firmware updates via <code>UpdateCapsule()</code> MAY return <code>EFI_UNSUPPORTED</code> , when called after the UEFI boot services have been exited.
	<i>See additional guidance.</i>

Chapter 6. BRS-I ACPI Requirements

The *Advanced Configuration and Power Interface Specification* provides the OS-centric view of system configuration, various hardware resources, events and power management.

This section defines the BRS-I mandatory and optional ACPI requirements on top of existing [3] and [9] specification requirements. Additional non-normative guidance may be found in the [firmware implementation guidance](#) section.



All content in this section is optional and recommended for BRS-B.

ID#	Requirement
ACPI_010	Be 64-bits clean. <ul style="list-style-type: none">• RSDT MUST NOT be implemented, with <code>RsdtdAddress</code> in RSDP set to 0.• 32-bit address fields MUST be 0.
See additional guidance.	
ACPI_020	Implement the hardware-reduced ACPI mode (no FACS table).
See additional guidance.	
ACPI_030	The Processor Properties Table (PPTT) MUST be implemented, even on systems with a simple hart topology.
ACPI_040	The PCI Memory-mapped Configuration Space (MCFG) table MUST NOT be present if it violates [18].
<i>Only compatible PCIe segments, exposed via ECAM (Enhanced Configuration Access Mechanism), may be described in the MCFG. The MCFG MUST NOT require vendor-specific OS support. See PCI Services ([3], Section 4) for more ACPI requirements relating to PCIe support. See additional guidance.</i>	
ACPI_050	A Serial Port Console Redirection Table [19] MUST be present on systems, where the graphics hardware is not present or not made available to an OS loader via the standard <code>UEFI EFI_GRAPHICS_OUTPUT_PROTOCOL</code> interface.
<i>In these cases, the table provides essential configuration for an early OS boot console.</i>	
ACPI_060	An SPCR table, if present, MUST meet the following requirements: <ul style="list-style-type: none">• Revision 4 or later of SPCR.• For NS16550-compatible UARTs:<ul style="list-style-type: none">◦ Use <code>Interface Type</code> 0x12 (16550-compatible with parameters defined in Generic Address Structure).◦ There MUST be a matching AML device object with compatible ID <code>RSCV0003</code>.
See additional guidance.	

ID#	Requirement
ACPI_080	Depending on the interrupt controller implemented by the system, PLIC or APLIC namespace devices MUST be present in the ACPI namespace along with MADT entries. See RVI ACPI IDs .
Also see AML_090 , AML_100 and additional guidance .	

6.1. BRS-I ACPI Methods and Objects

This section lists additional requirements for ACPI methods and objects.

[See additional guidance](#).

ID#	Requirement
AML_010	The Cache Coherency Attribute (_CCA) device method MUST be implemented.
<i>This object provides information about whether a device has to manage cache coherency and about hardware support. This object is mandatory for all devices that can access CPU-visible memory. ([3] Section 6.2.17).</i>	
AML_020	The Current Resource Setting (_CRS) device method for a PCIe Root Complex SHOULD NOT contain resources of type DWordIO , QWordIO or ExtendedIO .
<i>Legacy PCI I/O BARs are uncommon in modern PCIe devices and support for PCI I/O space may complicate configuration of PCIe Root Complex hardware in a compliant manner.</i>	
AML_030	The Possible Resource Settings (_PRS) and Set Resource Settings (_SRS) device method SHOULD NOT be implemented.
<i>ACPI resource descriptors are typically used to describe devices with fixed I/O regions that do not change. Flexible resource assignment is not supported by most modern ACPI OSes.</i>	
AML_040	Per-hart device objects MUST be defined under _SB (System Bus) namespace and not in the deprecated _PR (Processors) namespace.
AML_050	Systems supporting OS-directed hart performance control and power management MUST expose these via Collaborative Processor Performance Control (CPPC, [3] Section 8.4.6).
AML_060	Processor idle states MUST be described using Low Power Idle (_LPI , [3] Section 8.4.3).
AML_070	Systems with a Real-Time Clock on an OS-managed bus (e.g. I2C, subject to arbitration issues due to access to the bus by the OS) MUST implement the Time and Alarm Device (TAD).
Also see URT_020 .	
AML_080	Systems implementing a TAD MUST be functional without additional system-specific OS drivers.

ID#	Requirement
<i>In situations where the Time and Alarm Device (TAD) depends on a vendor-specific OS driver for correct function (SPI, I2C, etc), the TAD MUST be functional if the OS driver is not loaded. That is, when a dependent driver is loaded, an AML method switches further accesses to go through the driver-backed OperationRegion.</i>	
AML_090	PLIC and APLIC device objects MUST support the Global System Interrupt Base (_GSB , [3] Section 6.2.7) object. See additional guidance .
AML_100	Devices with Global System Interrupt (GSI, aka wired interrupt) resources MUST indicate the dependency on the corresponding interrupt controller using Operation Region Dependencies (_DEP , [3] Section 6.5.8). See additional guidance .
AML_110	UART device objects with ID RSCV0003 MUST implement Properties for UART Devices .

6.2. RVI-specific ACPI IDs

ACPI ID is used in the [_HID](#) (Hardware ID), [_CID](#) (Compatibility ID) or [_SUB](#) (Subsystem ID) objects as described in the ACPI Specification for devices, that do not have a standard enumeration mechanism. The ACPI ID consists of two parts: a vendor identifier followed by a product identifier.

Vendor IDs consist of 4 characters, each character being either an uppercase letter (A-Z) or a numeral (0-9). The vendor ID SHOULD be unique across the Industry and registered by the UEFI forum. For RVI standard devices, [RSCV](#) is the vendor ID registered. Vendor-specific devices can use an appropriate vendor ID registered for the manufacturer.

Product IDs are always four-character hexadecimal numbers (0-9 and A-F). The device manufacturer is responsible for assigning this identifier to each product model.

This document contains the canonical list of ACPI IDs for the namespace devices that adhere to the RVI specifications. The RVI task groups may make pull requests against this repository to request the allocation of ACPI ID for any new device.

ACPI ID	Device
RSCV0001	RISC-V Platform-Level Interrupt Controller (PLIC)
RSCV0002	RISC-V Advanced Platform-Level Interrupt Controller (APLIC)
RSCV0003	NS16650 UART compatible with an SPCR definition using Interface Type 0x12
<i>Also see Section 6.3.1.</i>	

6.3. RVI-specific ACPI Device Properties

This section defines the [_DSD](#) device properties [20] in the [rscv-](#) namespace.

Where explicit values are provided in a property definition, only these values must be used. System behavior with any other values is undefined.

Request for additional property names in the [rscv-](#) namespace should be made as a git pull request

to this document.

6.3.1. Properties for UART Devices

Property (<i>rscv-uart-*</i>)	Type	Description
<i>clock-frequency</i>	Integer	Clock feeding the IP block in Hz.
<i>A value of zero will preclude the ability to set the baud rate, or to configure a disabled device.</i>		
<i>reg-shift</i>	Integer	Quantity to shift the register offsets by.
<i>reg-io-width</i>	Integer	The size (in bytes) of the register accesses that should be performed on the device.
<i>1, 2, 4 or 8.</i>		
<i>rx-fifo-size</i>	Integer	The receive FIFO size (in bytes).

Chapter 7. BRS-I SMBIOS Requirements

The *System Management BIOS (SMBIOS) Reference Specification* defines a standard format for presenting management information about an implementation, mostly focusing on hardware components.

This section defines the BRS-I mandatory and optional SMBIOS requirements on top of existing [8] specification requirements. Additional non-normative guidance may be found in the [firmware implementation guidance](#) section.



All content in this section is optional and recommended for BRS-B.



The structures and fields in this section are defined in a manner consistent with the DMTF specification language ([8]).

ID#	Requirement
SMBIOS_010	A Baseboard/Module Information (Type 02) structure SHOULD be implemented.
SMBIOS_020	A System Enclosure/Chassis (Type 03) structure SHOULD be implemented.
<i>This relaxes the SMBIOS specification requirement.</i>	
SMBIOS_030	A Processor Information (Type 04) structure, meeting the additional Section 7.1 clarifications, MUST be implemented.
<i>This supersedes the RISC-V specific language in the SMBIOS specification ([8], Section 7.5.3.5).</i>	
SMBIOS_040	A Port Connector Information (Type 08) SHOULD be implemented.
SMBIOS_050	A System Slots (Type 09) structure MUST be implemented, when expansion slots are present.
SMBIOS_060	An OEM Strings (Type 11) structure SHOULD be implemented.
SMBIOS_070	A BIOS Language Information (Type 13) structure SHOULD be implemented.
SMBIOS_080	A Group Associations (Type 14) structure SHOULD be implemented.
SMBIOS_090	An IPMI Device Information (Type 38) structure MUST be implemented, when an IPMIv1.0 host interface is present.
SMBIOS_100	A System Power Supplies (Type 39) structure SHOULD be implemented.
SMBIOS_110	An Onboard Devices Extended Information (Type 41) structure SHOULD be implemented.
SMBIOS_120	A Redfish Host Interface (Type 42) structure MUST be implmented, when a Redfish host interface is present.
SMBIOS_130	A TPM Device (Type 43) structure MUST be implmented, when a TPM is present.
SMBIOS_140	A Processor Additional Information (Type 44) structure MUST be implemented.
<i>See the structure definitions below.</i>	
SMBIOS_150	A Firmware Inventory Information (Type 45) structure SHOULD be implemented.

7.1. Type 04 Processor Information



The information in this section supersedes the definitions in ([8], Section 7.5.3.4).

A processor is a grouping of harts in a physical package. In modern designs this MAY mean a SoC.

For RISC-V class CPUs, the **Processor ID** field contains two **DWORD**-formatted values describing the overall physical processor package vendor and version. For some implementations this may also be known as the SoC ID. The first **DWORD** (offsets 08h-0Bh) is the JEP-106 code for the vendor. The second **DWORD** (offsets 0Ch-0Fh) reflects vendor-specific part versioning.

For hart-specific vendor and revision information, please see Type 44 Processor-Specific Data structures.

7.2. Type 44 Processor-Specific Data

The processor-specific data structure fields are defined to follow the standard Processor-Specific Block fields ([8], Section 7.45.1).

The structure is valid for processors declared with **Processor Type** 07h (64-bit RISC-V) only.

A Type 44 structure needs to be provided for every hart meeting [Chapter 3](#) requirements.

Offset	Version	Name	Length	Value	Description
00h	0100h	Revision	WORD	Varies	See Section 7.3 .
02h	0100h	Hart ID	QWORD	Varies	The ID of this RISC-V hart.
0Ah	0100h	Machine Vendor ID	QWORD	Varies	The vendor ID of this RISC-V hart.
12h	0100h	Machine Architecture ID	QWORD	Varies	Base microarchitecture of the hart. Value of 0 is possible to indicate the field is not implemented. The combination of Machine Architecture ID and Machine Vendor ID should uniquely identify the type of hart microarchitecture that is implemented.
1Ah	0100h	Machine Implementation ID	QWORD	Varies	Unique encoding of the version of the processor implementation.

7.3. Processor-Specific Data Structure Versioning

The processor-specific data structure begins with a revision field to allow for future extensibility in a backwards-compatible manner.

The minor revision is to be incremented anytime new fields are added in a backwards-compatible manner. The major revision is to be incremented on backwards-incompatible changes.

Version	Bits 15:8+ Major revision	Bits 7:0+ Minor revision	Combined	Description
v1.0	01h	00h	0100h	First BRS-defined definition

Chapter 8. Firmware Implementation Guidance

The guidance section is non-normative, and covers certain implementation choices, suggestions, historical context, etc.

8.1. Recipes Guidance

8.1.1. BRS-I Recipe Guidance

Systems compliant to BRS-I can successfully boot an existing generic operating system image without system-specific customizations, yet this might result in an unoptimized experience and non-functioning I/O devices until further software updates are activated.

The best analogy would be a typical Intel Architecture motherboard from the early 2000s: you could install an OS on it, but the built-in graphics might be low-resolution and the sound, built-in network port or power management might not work out of the box. You could subsequently load the right drivers from the media coming with the board or fetch newest ones using a well-supported network adapter.

Requirements for heterogeneous performance harts (e.g. mix of "performance" and "efficiency" harts) are described in the "RISC-V Server Platform specification" [21]. The ACPI specification [3] has a chapter on "Collaborative Processor Performance Control".

8.2. UEFI Implementation Guidance

UEFI implementations run in 64-bit S-Mode, VS-Mode or HS-Mode, depending on whether virtualization is supported or used.

8.2.1. Privilege Levels

Different portions of system firmware might target a specific privilege level. In contrast, UEFI drivers, OS loaders and pre-boot applications need to operate in supervisor mode (either (V)S-Mode or HS-Mode), because they are UEFI-compliant executable images.

As an implementation choice, a UEFI firmware implementation may start execution in M-Mode. However it must switch to supervisor mode as part of initializing boot and runtime services for UEFI drivers and applications.

8.2.2. Firmware Update

`UpdateCapsule()` is only required before `ExitBootServices()` is called. The `UpdateCapsule()` implementation is expected to be suitable for use by generic firmware update services like fwupd and Windows Update. Both fwupd and Windows Update read the ESRT table to determine what firmware can be updated and use an EFI helper application to call `UpdateCapsule()` before `ExitBootServices()` is called.

8.2.3. PCIe

Every implementation of the `EFI_PCI_ROOT_BRIDGE_IO_PROTOCOL` provides the correct `Address Translation Offset` field to translate between the hart MMIO and bus addresses.

`EFI_PCI_ROOT_BRIDGE_IO_PROTOCOL_CONFIGURATION` structures report resources produced by the PCIe root bridges, not resources consumed by their register maps. In the cases where there are unpopulated PCIe slots behind a root bridge, `EFI_PCI_ROOT_BRIDGE_IO_PROTOCOL_CONFIGURATION` reports valid resources assigned (e.g. for hot plug), or reports no resources assigned.

8.3. ACPI Implementation Guidance

ACPI information is structured as tables with the address of the root of these tables known as Root System Description Table (RSDP) passed to the OS via a `EFI_ACPI_20_TABLE_GUID` configuration table in the UEFI firmware. The Operating System uses this address to locate all other ACPI tables.

Certain implementations may make use of the *RISC-V Functional Fixed Hardware Specification* [22].

8.3.1. 64-bits Clean

ACPI started as a specification for 32-bit systems, so certain tables with physical address pointers (e.g. RSDP, FADT) allow for reporting either 32-bit or 64-bit values using different fields. For the sake of simplicity and consistency, the BRS disallows the use 32-bit address fields in such structures and disallows the use of 32-bit only structures (thus, RSDT must not be implemented, as the XSDT is a direct replacement). Thus, the ACPI tables are allowed to be located in any part of the physical address space.

8.3.2. Hardware-Reduced ACPI

Compliant RISC-V systems only implement the hardware-reduced ACPI model [3] (Section 4.1). This means the hardware portion of [3] (Section 4) is not required or supported. All functionality is instead provided through equivalent software-defined interfaces and the complexity in supporting ACPI is reduced.

8.3.3. Table Guidance

[Table 4](#) summarizes the minimum set of structures that must exist to support basic booting of RISC-V system with ACPI support. [Table 5](#) lists additional possible ACPI tables based on the optional features that can be supported. The latter is not meant to be exhaustive and mostly focuses on tables that have specific guidance or that are expected to be frequently implemented.

Table 4. Minimum required ACPI System Description Tables

ACPI Table	ACPI Section	Note
Root System Description Pointer (RSDP)	5.2.5	See high-level requirements .
Extended System Description Table (XSDT)	5.2.8	Contains pointers to other tables.

ACPI Table	ACPI Section	Note
Fixed ACPI Description Table (FADT)	5.2.9	See ACPI_020 , Section 8.3.2 and the notes below .
Differentiated System Description Table (DSDT)	5.2.11.1	See Section 6.1 and the notes below .
Multiple APIC Description Table (MADT)	5.2.12	See the notes below
RISC-V Hart Capabilities Table (RHCT)	New	Communicates information about certain capabilities like ISA string, cache and MMU info.
Processor Properties Topology Table (PPTT)	5.2.29	See ACPI_030

Table 5. Additional ACPI System Description Tables based on feature support.

ACPI Table	ACPI Section	Note
Memory-mapped Configuration space (MCFG)	[18]	See ACPI_040 and the notes below
Secondary System Description Table (SSDT)	5.2.11.2	See Section 6.1 and the notes below .
Serial Port Console Redirection (SPCR)	[19]	See ACPI_060 and the notes below
ACPI Table for TPM 2.0 (TPM2)	[23]	If the system supports TPM 2.0
System Resource Affinity Table (SRAT)	5.2.16	If the system supports NUMA
System Locality Information Table (SLIT)	5.2.17	If the system supports NUMA
Boot Error Record Table (BERT)	18.3.1	If APEI is supported
Error Injection Table (EINJ)	18.6.1	If APEI is supported
Error Record Serialization Table (ERST)	18.5	If APEI is supported
Hardware Error Source Table (HEST)	18.3.2	If APEI is supported

8.3.4. DSDT and SSDTs

The ACPI name space describes devices which cannot be enumerated by any other standard ways. These typically include SoC embedded memory-mapped I/O devices, such as UARTs, [PCIe](#) or CXL root complexes, GPIO controllers, etc.

It's an implementation choice if the ACPI name space is defined solely with a DSDT or with any additional SSDTs. For example, a UEFI implementation may choose to use SSDTs to:

- describe devices that vary across SoC SKUs, revisions or variants.
- describe devices, where the backing AML is generated or patched at boot time.

8.3.5. FADT

[3] (Section 5.2.9) provides guidance on filling the Fixed ACPI Description Table for HW-reduced ACPI.

Don't forget to select an appropriate Preferred PM Profile.

8.3.6. MADT

RINTC (per-hart) structures are mandatory. Depending on the interrupt controller implemented by the system, the MADT will also contain either PLIC or IMSIC/APLIC structures.

Entry ordering can be correlated with initialization order by an OS, but should not be taken to reflect affinity in resource sharing, e.g. sockets, caches, etc. RINTC hart ID and ACPI Processor UID should not be decoded in a system-specific manner to divine CPU topology. The PPTT Processor Properties Topology Table (PPTT) is to be used to describe affinities.

8.3.7. PLIC/APLIC Namespace Devices

Here's an example of an ASL excerpt satisfying [ACPI_080](#), [AML_090](#) and [AML_100](#) requirements.

```
Scope (\_SB)
{
    Device (IC00)
    {
        Name (_HID, "RSCV0001") // _HID: Hardware ID
        Name (_UID, Zero) // _UID: Unique ID
        Method(_GSB) {
            Return (0x10) // Global System Interrupt Base for this PLIC starts at 16
        }
    }
    Device (DEV1)
    {
        ...
        Name (_DEP, Package() {\_SB.IC00})
        Name (_CRS, ResourceTemplate () // _CRS: Current Resource Settings
        {
            Memory32Fixed (ReadWrite,
                0x10010000, // Address Base.
                0x00010000, // Address Length
            )
            Interrupt (ResourceConsumer, Level, ActiveHigh, Exclusive, ,,)
            {
                0x10,
            }
        })
    }
}
```

```
}  
}
```

8.3.8. PCIe

On some architectures, it became an industry accepted norm to describe PCIe implementations not compliant to the *PCI Firmware Specification* [18] using specification-defined ACPI tables and objects. RISC-V systems compliant to the BRS must only expose ECAM-compatible implementations using the MCFG and the standard AML Hardware ID (`_HID`) `PNP0A08` and Compatible ID (`_CID`) `PNP0A03`, and must not rely on ACPI table header information or other out-of-band means of detecting quirked behavior.

Some minor incompatibilities, such as incorrect CFG0 filtering, broken BARs/capabilities for RCs, embedded switches/bridges or embedded endpoints can be handled by emulating ECAM accesses in privileged firmware (e.g. M-mode) or similar facilities (e.g. a hypervisor).

Non-compliant implementations must be exposed using vendor-specific mechanisms (e.g. AML object with custom `_HID`, custom vendor-specific ACPI table if necessary). In cases where such PCIe implementations are only used to expose a fixed non-removable device (e.g. USB host controller or NVMe), the device could be exposed via a DSDT/SSDT MMIO device object without making the OS aware of the underlying PCIe connection.

8.3.9. SPCR

Early serial console can be implemented using either an NS16550 UART (SPCR `Interface Type` 0x12) or SBI console (SPCR `Interface Type` 0x15). When SPCR describes SBI console, the OS must use the SBI Probe extension (`FID #3`) to detect the appropriate facilities, e.g. the Debug Console Extension (`DBCN`) or the deprecated legacy console EIDs.

The new `Precise Baud Rate` field, introduced in [19] rev. 4, allows describing rates faster than 115200 baud for NS16550-compatible UARTS.

Hardware not capable of interrupt-driven operation and SBI console should be described with `Interrupt Type` of 0 and `Global System Interrupt` of 0.

8.4. SMBIOS Implementation Guidance

Note the DMTF requirements on the 64-bit SMBIOS 3.0 entry point ([8] Section 5.2.2) and data structures ([8] Section 6.2).

8.4.1. Type 44 Processor-Specific Data

The `Machine Vendor ID`, `Machine Architecture ID`, and `Machine Implementation ID` fields typically reflect the `mvendorid`, `marchid` and `mimpid` CSRs respectively.

Bibliography

- [1] “Key words for use in RFCs to Indicate Requirement Levels.” [Online]. Available: datatracker.ietf.org/doc/html/rfc2119.
- [2] “PCI Express® Base Specification Revision 6.0.” [Online]. Available: pcisig.com/pci-express-6.0-specification.
- [3] “Advanced Configuration and Power Interface Specification 6.6.” [Online]. Available: uefi.org/specifications.
- [4] “DeviceTree.” [Online]. Available: www.devicetree.org/.
- [5] “Embedded Base Boot Requirements Specification 2.1.0.” [Online]. Available: github.com/ARM-software/ebbr/releases/download/v2.1.0/ebbr-v2.1.0.pdf.
- [6] “RISC-V Profile.” [Online]. Available: github.com/riscv/riscv-profiles.
- [7] “RISC-V Supervisor Binary Interface Specification.” [Online]. Available: github.com/riscv-non-isa/riscv-sbi-doc.
- [8] “System Management BIOS (SMBIOS) Reference Specification 3.7.0.” [Online]. Available: www.dmtf.org/standards/smbios.
- [9] “Unified Extensible Firmware Interface Specification 2.11.” [Online]. Available: uefi.org/specifications.
- [10] “RISC-V ‘stimecmp / vstimecmp’ Extension.” 2021, [Online]. Available: github.com/riscv/riscv-time-compare.
- [11] “The RISC-V Advanced Interrupt Architecture.” 2023, [Online]. Available: github.com/riscv/riscv-ai.
- [12] “RISC-V Indirect CSR Access (Smcsrind/Sscsrind).” 2023, [Online]. Available: github.com/riscv/riscv-indirect-csr-access.
- [13] “RISC-V Supervisor Counter Delegation Specification (Smcdeleg/Ssccfg).” 2024, [Online]. Available: github.com/riscv/riscv-smcdeleg-ssccfg.
- [14] “SIG: Performance Analysis.” 2024, [Online]. Available: lists.riscv.org/g/sig-perf-analysis.
- [15] “UEFI memory mitigations.” [Online]. Available: learn.microsoft.com/en-us/windows-hardware/drivers/bringup/uefi-ca-memory-mitigation-requirements.
- [16] “UEFI Platform Initialization Specification 1.9.” [Online]. Available: uefi.org/specifications.
- [17] “TCG EFI Platform Specification.” [Online]. Available: trustedcomputinggroup.org/resource/tcg-efi-platform-specification/.
- [18] “PCI Firmware Specification Revision 3.3.” [Online]. Available: members.pcisig.com/wg/PCI-SIG/document/folder/862.

- [19] “Serial Port Console Redirection Table (SPCR).” [Online]. Available: learn.microsoft.com/en-us/windows-hardware/drivers/serports/serial-port-console-redirection-table.
- [20] “_DSD (Device Specific Data) Implementation Guide.” [Online]. Available: github.com/UEFI/DSD-Guide/blob/main/dsd-guide.pdf.
- [21] “RISC-V Server Platform Specification.” [Online]. Available: github.com/riscv-non-isa/riscv-server-platform.
- [22] “RISC-V Functional Fixed Hardware Specification.” [Online]. Available: github.com/riscv-non-isa/riscv-acpi-ffh.
- [23] “TCG ACPI Specification.” [Online]. Available: trustedcomputinggroup.org/resource/tcg-acpi-specification/.