

The next generation QEMU functional testing framework



Daniel P. Berrangé <berrange@redhat.com>

Thomas Huth <thuth@redhat.com>



a friendly emu that holds a test tube with some chemicals that just exploded and covered the emu in coal dust

Legal

- Disclaimer: Opinions are our own and not necessarily the views of our employer
- The somewhat absurd emu pictures have been generated with AI (Gemini; prompts noted inline). All other content in this presentation is the work of humans.



a cartoon emu as a solicitor in court



comic-style picture of a friendly and smiling emu that flies through the air using rockets



A little bit of history

- QEMU had many test suites (unit tests, “qtests”, “iotests”, TCG-tests, ...)
- Missing end-to-end functional testing of VMs with real payloads (kernel, etc.)
- Needed to write tests more easily. e.g. python + helpers for assets & caching
- Maintainers from the Avocado project offered help in 2018
- The Avocado-based functional test suite was introduced (initially called “acceptance” tests, later renamed to “avocado” tests)





The issues with Avocado for QEMU



- Avocado is a fairly complex system
- The average QEMU developer is a C hacker, probably not a Python expert
- If something did not work, QEMU developers rarely understood how to fix it
- Little proactive help from the Avocado developers after the initial code drop
 - But no other QEMU maintainers felt responsible for the subsystem
- Attempts to upgrade Avocado to a newer release failed for various reasons
 - Stuck with outdated & unmaintained Avocado v88.1 (from 2021)
- Last straw in 2024: Avocado v88.1 not compatible with Python 3.12
 - A solution had to be found ASAP

(ok, this picture was rather created with gimp from the previous one, since the AI was not able to come up with a nice picture of an emu putting its head into the ground)





Re-considering the testing situation



- Avocado was performing two jobs, providing
 - (a) test runner / harness
 - (b) test creation infrastructure APIs
- 2018: A zoo of historically grown test harness for various test subsystems
- 2024: Meson provides the parallel test runner for most QEMU frameworks
 - Avocado tests were an oddball here (v88 was single-threaded)
 - Avocado test harness results difficult to interpret & complicated debugging failures
- 2022: “Introduce new acpi/smbios python tests using biosbits”
 - Tried a new stand-alone Python-based functional test separate from Avocado
- Could Avocado tests be migrated to standalone python tests run by Meson ?





Example of an Avocado test (slightly modified)

```
from avocado_gemu import QemuSystemTest, wait_for_console_pattern
from avocado.utils import archive
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    timeout = 90
```

```
    @skipUnless(os.getenv('QEMU_TEST_FLAKY_TESTS'), 'Test might be unstable')
```

```
    def test_arm_canona1100(self):
```

```
        """
```

```
        :avocado: tags=arch:arm
```

```
        :avocado: tags=machine:canon-a1100
```

```
        """
```

```
        tar_url = ('https://gemu-advcal.gitlab.io/.../day18.tar.xz')
```

```
        tar_hash = '068b5fc4242b29381acee94713509f8a876e9db6'
```

```
        file_path = self.fetch_asset(tar_url, asset_hash=tar_hash)
```

```
        archive.extract(file_path, self.workdir)
```

```
        self.vm.add_args('-bios', self.workdir + '/day18/barebox.canon-a1100.bin')
```

```
        self.vm.set_console()
```

```
        self.vm.launch()
```

```
        wait_for_console_pattern(self, 'running /env/bin/init')
```



a cartoon emu viewing the matrix



Example of an Avocado test (slightly modified)

```
from avocado_qemu import QemuSystemTest, wait_for_console_pattern
from avocado.utils import archive
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    timeout = 90
```

← **Basic test class**

```
@skipUnless(os.getenv('QEMU_TEST_FLAKY_TESTS'), 'Test might be unstable')
```

```
def test_arm_canona1100(self):
```

```
    """
```

```
    :avocado: tags=arch:arm
```

```
    :avocado: tags=machine:canon-a1100
```

```
    """
```

```
    tar_url = ('https://qemu-advcal.gitlab.io/.../day18.tar.xz')
```

```
    tar_hash = '068b5fc4242b29381acee94713509f8a876e9db6'
```

```
    file_path = self.fetch_asset(tar_url, asset_hash=tar_hash)
```

```
    archive.extract(file_path, self.workdir)
```

```
    self.vm.add_args('-bios', self.workdir + '/day18/barebox.canon-a1100.bin')
```

```
    self.vm.set_console()
```

```
    self.vm.launch()
```

```
    wait_for_console_pattern(self, 'running /env/bin/init')
```





Example of an Avocado test (slightly modified)

```
from avocado_gemu import QemuSystemTest, wait_for_console_pattern
from avocado.utils import archive
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    timeout = 90
```

← **Timeout setting for the test runner**

```
@skipUnless(os.getenv(QEMU_TEST_FLAKY_TESTS), 'Test might be unstable')
```

```
def test_arm_canona1100(self):
```

← **Subtest method**

```
    """
```

```
    :avocado: tags=arch:arm
```

```
    :avocado: tags=machine:canon-a1100
```

```
    """
```

```
    tar_url = ('https://gemu-advcal.gitlab.io/.../day18.tar.xz')
```

```
    tar_hash = '068b5fc4242b29381acee94713509f8a876e9db6'
```

```
    file_path = self.fetch_asset(tar_url, asset_hash=tar_hash)
```

```
    archive.extract(file_path, self.workdir)
```

```
    self.vm.add_args('-bios', self.workdir + '/day18/barebox.canon-a1100.bin')
```

```
    self.vm.set_console()
```

```
    self.vm.launch()
```

```
    wait_for_console_pattern(self, 'running /env/bin/init')
```





Example of an Avocado test (slightly modified)

```
from avocado_qemu import QemuSystemTest, wait_for_console_pattern
from avocado.utils import archive
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    timeout = 90
```

Decorator

```
    @skipUnless(os.getenv(QEMU_TEST_FLAKY_TESTS), 'Test might be unstable')
```

```
    def test_arm_canona1100(self):
```

```
        """
```

```
        :avocado: tags=arch:arm
```

```
        :avocado: tags=machine:canon-a1100
```

```
        """
```

```
        tar_url = ('https://qemu-advcal.gitlab.io/.../day18.tar.xz')
```

```
        tar_hash = '068b5fc4242b29381acee94713509f8a876e9db6'
```

```
        file_path = self.fetch_asset(tar_url, asset_hash=tar_hash)
```

```
        archive.extract(file_path, self.workdir)
```

```
        self.vm.add_args('-bios', self.workdir + '/day18/barebox.canon-a1100.bin')
```

```
        self.vm.set_console()
```

```
        self.vm.launch()
```

```
        wait_for_console_pattern(self, 'running /env/bin/init')
```





Example of an Avocado test (slightly modified)

```
from avocado_gemu import QemuSystemTest, wait_for_console_pattern
from avocado.utils import archive
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    timeout = 90
```

```
    @skipUnless(os.getenv('QEMU_TEST_FLAKY_TESTS'), 'Test might be unstable')
```

```
    def test_arm_canona1100(self):
```

```
        """
```

```
        :avocado: tags=arch:arm
```

```
        :avocado: tags=machine:canon-a1100
```

```
        """
```

```
        tar_url = ('https://gemu-advcal.gitlab.io/.../day18.tar.xz')
```

```
        tar_hash = '068b5fc4242b29381acee94713509f8a876e9db6'
```

```
        file_path = self.fetch_asset(tar_url, asset_hash=tar_hash)
```

```
        archive.extract(file_path, self.workdir)
```

```
        self.vm.add_args('-bios', self.workdir + '/day18/barebox.canon-a1100.bin')
```

```
        self.vm.set_console()
```

```
        self.vm.launch()
```

```
        wait_for_console_pattern(self, 'running /env/bin/init')
```



← **Tags**



Example of an Avocado test (slightly modified)

```
from avocado_gemu import QemuSystemTest, wait_for_console_pattern
from avocado.utils import archive
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    timeout = 90
```

```
    @skipUnless(os.getenv('QEMU_TEST_FLAKY_TESTS'), 'Test might be unstable')
```

```
    def test_arm_canona1100(self):
```

```
        """
```

```
        :avocado: tags=arch:arm
```

```
        :avocado: tags=machine:canon-a1100
```

```
        """
```

```
        tar_url = ('https://gemu-advcal.gitlab.io/.../day18.tar.xz')
```

```
        tar_hash = '068b5fc4242b29381acee94713509f8a876e9db6'
```

```
        file_path = self.fetch_asset(tar_url, asset_hash=tar_hash)
```

```
        archive.extract(file_path, self.workdir)
```

```
        self.vm.add_args('-bios', self.workdir + '/day18/barebox.canon-a1100.bin')
```

```
        self.vm.set_console()
```

```
        self.vm.launch()
```

```
        wait_for_console_pattern(self, 'running /env/bin/init')
```

Asset handling





Requirements for a new test framework

What?	Replacement
Test runner	Meson test runner (+ pycotap)
Test discovery	meson.build + unittest class (+ pycotap)
Basic test class	Recycle glue code (based on unittest class)
Decorators	Python unittest class (+ adding our own)
Asset download, caching & extracting	Replace with our own implementation
Logging	Python logging + custom setup
Tags (for selecting subsets of tests)	No replacement (possibly decorators?)



a cartoon emu writing a list of parts





Test files discovery and timeouts



- Avocado
 - Scanned all *.py files on invocation for available tests
 - A “timeout” variable within the test class for test timeout
- New functional tests
 - .py files and timeouts listed in tests/functional/*/meson.build files
- Slightly more work to add a test...
 - ... but it's the common pattern with meson to explicitly list files
- Meson allows to specify a priority
 - Meson can run all subsystem tests in parallel (qtests, iotests, unit, ...)
 - High priority for tests with large timeout makes long tests start first





How to handle subtests with meson



a cartoon emu surrounded by lots of baby emus

- Python ‘unittest’ class can be used for discovering & running subtest methods within a file. For running test standalone, simply add:

```
if __name__ == '__main__':  
    QemuSystemTest.main()
```


- By default, the meson test runner treats one file as one big test
- For showing progress on subtest level: TAP – Test anything protocol
- Tried a bunch of implementations (‘tappy’, ...), but none worked very well with the meson test runner (issues with stderr)
- Finally found ‘Pycotap’ that does the job and is very small!
 - Shipped as a wheel with QEMU now





Handle large test assets

- Most tests use assets (kernel, initrd, rootfs, firmware, etc.)
- Downloading assets is potentially quite slow
 - could cause test timeouts if done on demand
- Avocado had a local cache of assets that must be retained
- Asset class instances declared as class level variables
- Launching test with 'QEMU_TEST_PRECACHE=1' pre-caches asset files
- If any assets fail to download, CI jobs will skip the affected tests
 - ...except HTTP 404 codes, which are likely non-transient errors
- Added locking/waiting logic for downloading assets in parallel



a cartoon emu picking penguins from a tree





File management



a cartoon emu in a library covered in falling books

- Tests need to reference & create files in various locations
- Helper APIs provided on QemuBaseTest to standardize locations
 - `self.socket_dir()` => location for UNIX sockets
 - `self.data_file(...)` => source file relative to tests/functional dir
 - `self.build_file(...)` => build system created file relative to root of build dir
 - `self.scratch_file(...)` => any file created by test case, auto-deleted on exit
 - `self.log_file(...)` => any file recording log messages, uploaded as asset in CI jobs
 - `self.plugin_file(...)` => any TCG plugin relative to tests/tcg/plugins/
- APIs construct paths from components
 - `qualified_path = self.scratch_file("foo", "bar", "wizz")`
- Not
 - `qualified_path = self.scratch_file("foo/bar/wizz")`





Archive management

- Many assets are compressed or archives
- Lots of different python APIs for each format
- Helper APIs provided on QemuBaseTest to give easy access
- Archive extraction (tar, zip, cpio, deb)
 - `self.archive_extract(self.ASSET_.....,)` => unpacks to 'scratch_file' location
- File decompression (gz, xz, zstd)
 - `self.uncompress(self.ASSET_...,)` => uncompressed to 'scratch_file' location
- Formats guessed from archive URL path file extension



a cartoon emu wheeling a crate away from the viewer in a warehouse with crates as far as the eye can see





Decorators



- Typical python testing practice is to use decorators to control execution
- Functional test system provides standard decorators
 - `skipIfMissingCommands` => check 'binary' in \$PATH
 - `skipIfOperatingSystem` => exclude listed host OS
 - `skipIfNotMachine` => require a specific VM machine type
 - `skipFlakyTest` => don't run non-deterministic tests (GitLab issue URL required)
 - `skipUntrustedTest` => don't run potentially dangerous tests
 - `skipBigDataTest` => don't run tests which create huge files (> ~500 MB)
 - `skipSlowTest` => don't run tests which are excessively slow (many minutes)
 - `skipIfMissingImports` => check 'module' in \$PYTHONPATH
 - `skipLockedMemoryTest` => require permission to lock RAM





Example of a new functional test

```
from qemu_test import QemuSystemTest, Asset, skipFlakyTest
from qemu_test import wait_for_console_pattern
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    ASSET_BIOS = Asset('https://qemu-advcal.gitlab.io/.../day18.tar.xz',
                       '28e71874ce985be66b7fd1345ed88cb2523b982f899c8d2900d6353054a1be49')
```

```
    @skipFlakyTest('https://gitlab.com/qemu-project/qemu/-/issues/xyz')
```

```
    def test_arm_canona1100(self):
        self.set_machine('canon-a1100')
```

```
        bios = self.archive_extract(self.ASSET_BIOS,
                                     member="day18/barebox.canon-a1100.bin")
```

```
        self.vm.set_console()
```

```
        self.vm.add_args('-bios', bios)
```

```
        self.vm.launch()
```

```
        wait_for_console_pattern(self, 'running /env/bin/init')
```

```
if __name__ == '__main__':
    QemuSystemTest.main()
```



a cartoon emu inside the matrix dodging flying penguins





Example of a new functional test

```
from qemu_test import QemuSystemTest, Asset, skipFlakyTest
from qemu_test import wait_for_console_pattern
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    ASSET_BIOS = Asset('https://qemu-advcal.gitlab.io/.../day18.tar.xz',
                       '28e71874ce985be66b7fd1345ed88cb2523b982f899c8d2900d6353054a1be49')
```

```
    @skipFlakyTest('https://gitlab.com/qemu-project/qemu/-/issues/xyz')
    def test_arm_canona1100(self):
```

```
        self.set_machine('canon-a1100')
```

```
        bios = self.archive_extract(self.ASSET_BIOS,
                                     member="day18/barebox.canon-a1100.bin")
```

```
        self.vm.set_console()
```

```
        self.vm.add_args('-bios', bios)
```

```
        self.vm.launch()
```

```
        wait_for_console_pattern(self, 'running /env/bin/init')
```

```
if __name__ == '__main__':
    QemuSystemTest.main()
```

Replacement
for the tags



a cartoon emu inside the matrix dodging flying penguins





Example of a new functional test

```
from qemu_test import QemuSystemTest, Asset, skipFlakyTest
from qemu_test import wait_for_console_pattern
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    ASSET_BIOS = Asset('https://qemu-advcal.gitlab.io/.../day18.tar.xz',
                       '28e71874ce985be66b7fd1345ed88cb2523b982f899c8d2900d6353054a1be49')
```

```
    @skipFlakyTest('https://gitlab.com/qemu-project/qemu/-/issues/xyz')
    def test_arm_canona1100(self):
        self.set_machine('canon-a1100')
```

New
asset
handling

```
        bios = self.archive_extract(self.ASSET_BIOS,
                                     member="day18/barebox.canon-a1100.bin")
```

```
        self.vm.set_console()
        self.vm.add_args('-bios', bios)
        self.vm.launch()
        wait_for_console_pattern(self, 'running /env/bin/init')
```

```
if __name__ == '__main__':
    QemuSystemTest.main()
```



a cartoon emu inside the matrix dodging flying penguins





Example of a new functional test

```
from qemu_test import QemuSystemTest, Asset, skipFlakyTest
from qemu_test import wait_for_console_pattern
```

```
class CanonA1100Machine(QemuSystemTest):
```

```
    ASSET_BIOS = Asset('https://qemu-advcal.gitlab.io/.../day18.tar.xz',
                       '28e71874ce985be66b7fd1345ed88cb2523b982f899c8d2900d6353054a1be49')
```

```
    @skipFlakyTest('https://gitlab.com/qemu-project/qemu/-/issues/xyz')
```

```
    def test_arm_canona1100(self):
        self.set_machine('canon-a1100')
```

```
        bios = self.archive_extract(self.ASSET_BIOS,
                                     member="day18/barebox.canon-a1100.bin")
```

```
        self.vm.set_console()
```

```
        self.vm.add_args('-bios', bios)
```

```
        self.vm.launch()
```

```
        wait_for_console_pattern(self, 'running /env/bin/init')
```

```
if __name__ == '__main__':
    QemuSystemTest.main()
```

For running
standalone



a cartoon emu inside the matrix dodging flying penguins





Troubleshooting



a cartoon emu pulling apart computers in a server room tangled in cables


- Run tests standalone outside meson for easier debugging (strace)
 - `$ export PYTHONPATH=../python:../tests/functional`
 - `$ export QEMU_TEST_QEMU_BINARY=$PWD/qemu-system-x86_64`
 - `$ pyvenv/bin/python3 ../tests/functional/test_file.py`
- Getting information out of CI infra is always a challenge
- Logging is critical to understanding failures
 - `$BUILD/tests/functional/x86_64/$TEST_FILE.$TEST_CLASS.$TEST_METHOD/`
 - `base.log` ⇒ python 'logging' output from test class (includes CLI args of QEMU)
 - `console.log` ⇒ serial console output from QEMU guest
 - `default.log` ⇒ stdout/err from spawned QEMU
- All log files are uploaded as artifacts in GitLab CI jobs





Integration into the test suites

- “make check-functional” (or “make check-functional-ppc” etc.)
- Don’t want to do this by default during a normal “make check”
(a good internet connection is required for downloading the assets!)
- But some tests don’t need assets, i.e. should be run by default
 - Need a way to distinguish them
- Existing test suites already have speed classes:
 - *quick*, *slow* and *thorough*
- Only *quick* tests are run by default, i.e. add tests without assets here
- Add functional tests with assets to the *thorough* category
 - `make -j$(nproc) check SPEED=thorough`



a cartoon emu working on an assembly line in a computer factory





Demo



```
thuth@thuth-p1g4:~/tmp/qemu-build

$ make -j$(nproc) check-functional
[0/1] Running external command precache-functional (wrapped by meson to set env)
make[1]: Entering directory '/home/thuth/tmp/qemu-build'
[1/36] Generating qemu-version.h with a custom command (wrapped by meson to capture output)
/home/thuth/tmp/qemu-build/pyvenv/bin/meson test --no-rebuild -t 1 --setup thorough --num-processes 16 --print-errorlogs --suite func --suite func-quick --suite func-thorough
1/267 qemu:func-thorough+func-ppc64-thorough+thorough / func-ppc64-hv SKIP 0.13s 0 subtests passed
2/267 qemu:func-thorough+func-aarch64-thorough+thorough / func-aarch64-smmu SKIP 9.03s 0 subtests passed
3/267 qemu:func-thorough+func-aarch64-thorough+thorough / func-aarch64-raspi4 OK 39.32s 2 subtests passed
4/267 qemu:func-thorough+func-arm-thorough+thorough / func-arm-aspeed_gb200nv1_bmc OK 41.92s 1 subtests passed
5/267 qemu:func-thorough+func-mipsel-thorough+thorough / func-mipsel-replay SKIP 0.32s 0 subtests passed
6/267 qemu:func-thorough+func-arm-thorough+thorough / func-arm-aspeed_catalina OK 101.70s 1 subtests passed
7/267 qemu:func-thorough+func-aarch64-thorough+thorough / func-aarch64-device_passthrough OK 104.67s 1 subtests passed
8/267 qemu:func-thorough+func-arm-thorough+thorough / func-arm-orangepi OK 110.46s 3 subtests passed
9/267 qemu:func-thorough+func-arm-thorough+thorough / func-arm-aspeed_bletchley OK 113.38s 1 subtests passed
10/267 qemu:func-thorough+func-arm-thorough+thorough / func-arm-bpim2u OK 118.96s 3 subtests passed
11/267 qemu:func-thorough+func-aarch64-thorough+thorough / func-aarch64-virt_gpu OK 136.45s 2 subtests passed
12/267 qemu:func-thorough+func-x86_64-thorough+thorough / func-x86_64-replay OK 60.75s 1 subtests passed
[13-28/267] qemu:func-thorough+func-arm-thorough+thorough / func-arm-aspeed_ast2500 174/720.0s 1 subtests passed
```

a cartoon emu presenting a new computer product on stage for the "Emu Tools" company launch event





Future plans

- Evicting obsolete assets from the download cache
- Improve error handling, e.g. if QEMU crashes
- Enforce mypy, flake8, pylint, etc; format with 'black'
- Add more test for uncovered areas
 - If you know how to test one of the missing machines, please help:

<https://wiki.qemu.org/Testing/Machines>



a cartoon emu piloting a flying saucer across the galaxy





Any questions ?

(Or find us in the hallway later)